



BRENT OZAR
UNLIMITED®

Why Defragmenting Your Indexes Isn't Helping



BRENT OZAR
UNLIMITED

99-05: dev, architect, DBA
05-08: DBA, VM, SAN admin
08-10: MCM, Quest Software
Since: consulting DBA

www.BrentOzar.com
Help@BrentOzar.com

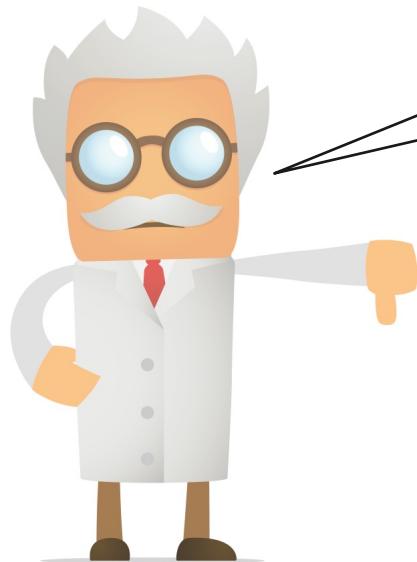


You're here because
performance sucks.



Success means
a performance boost
that end users notice.





"Fragmentation sucks."

The two kinds of fragmentation

Why the fix makes things worse

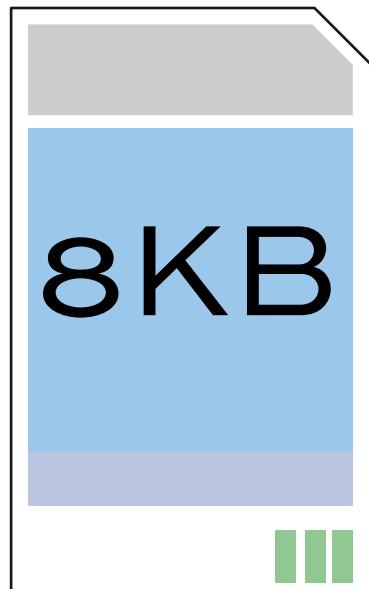
The better performance fix

When people insist on rebuilds

What to do overnight instead



**Data is
stored in
8KB pages.**

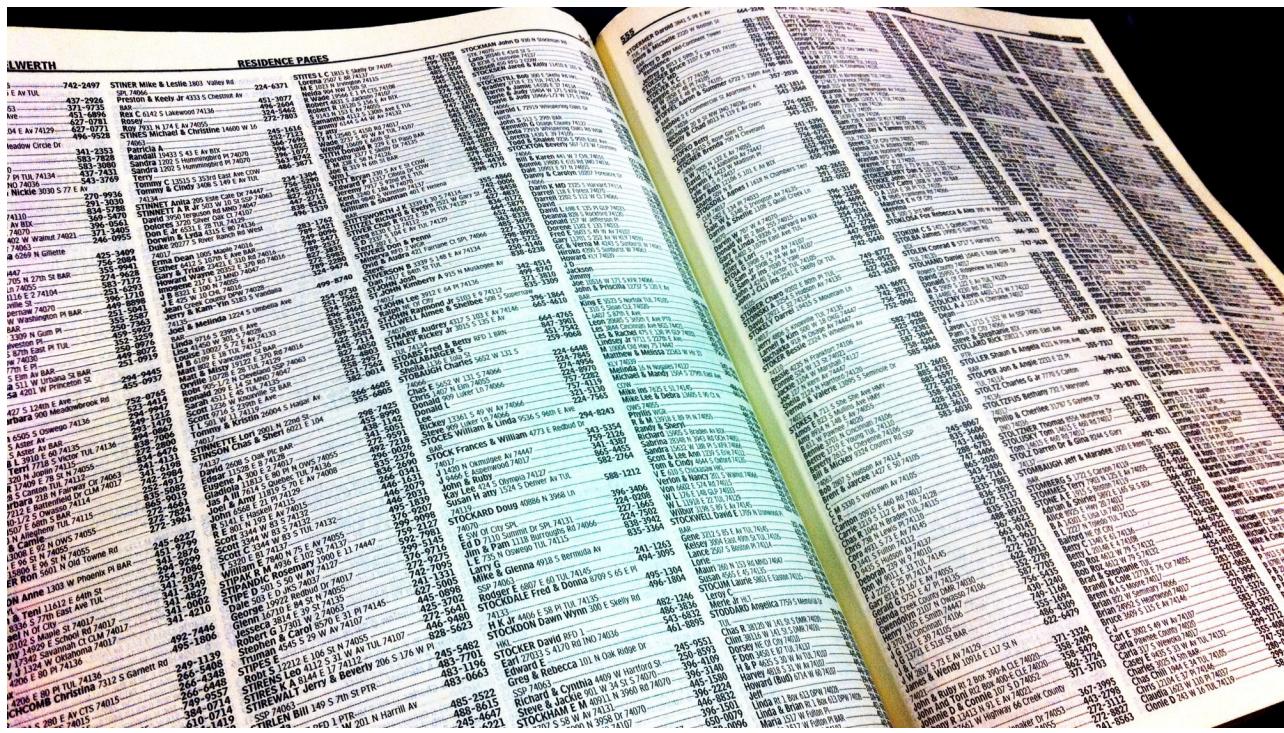


Page Header

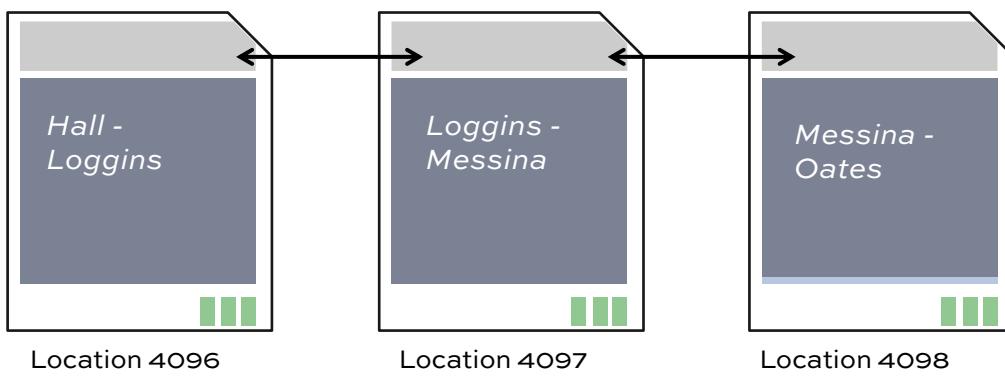
Index **OR**
Data Rows

Slot Array

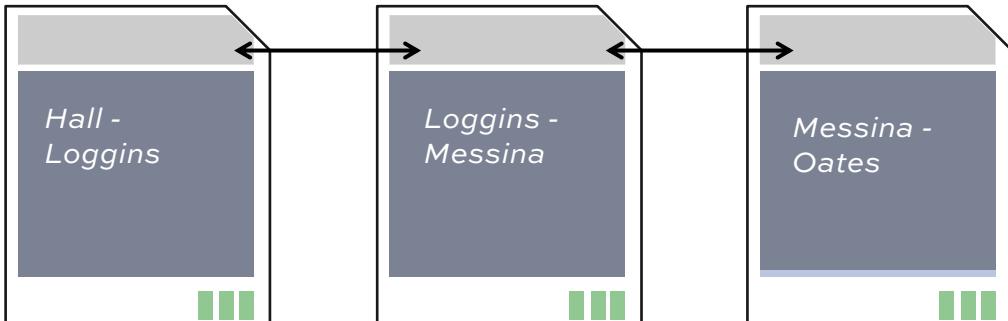




In a freshly rebuilt index, the 8K pages are full,
and let's say they're all in order on disk.



McDonald, Michael moves to our town.

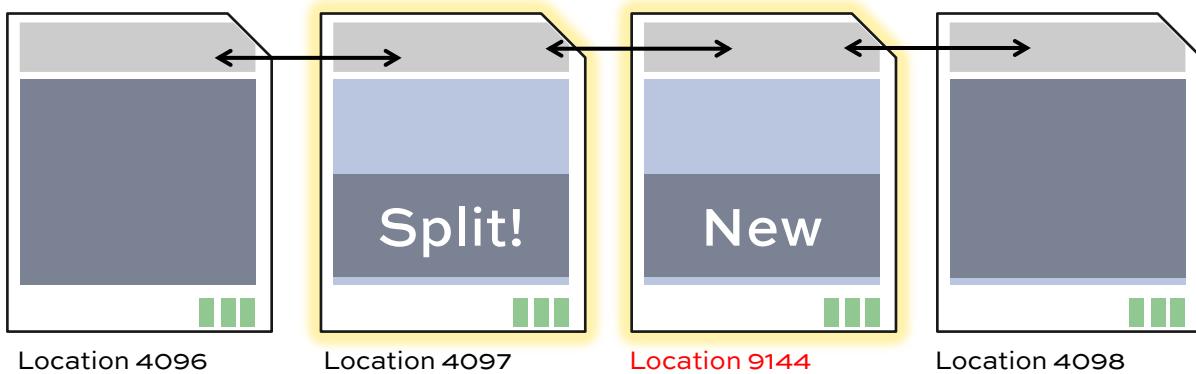


He needs to go
here, but there's
no empty space.



Page split!

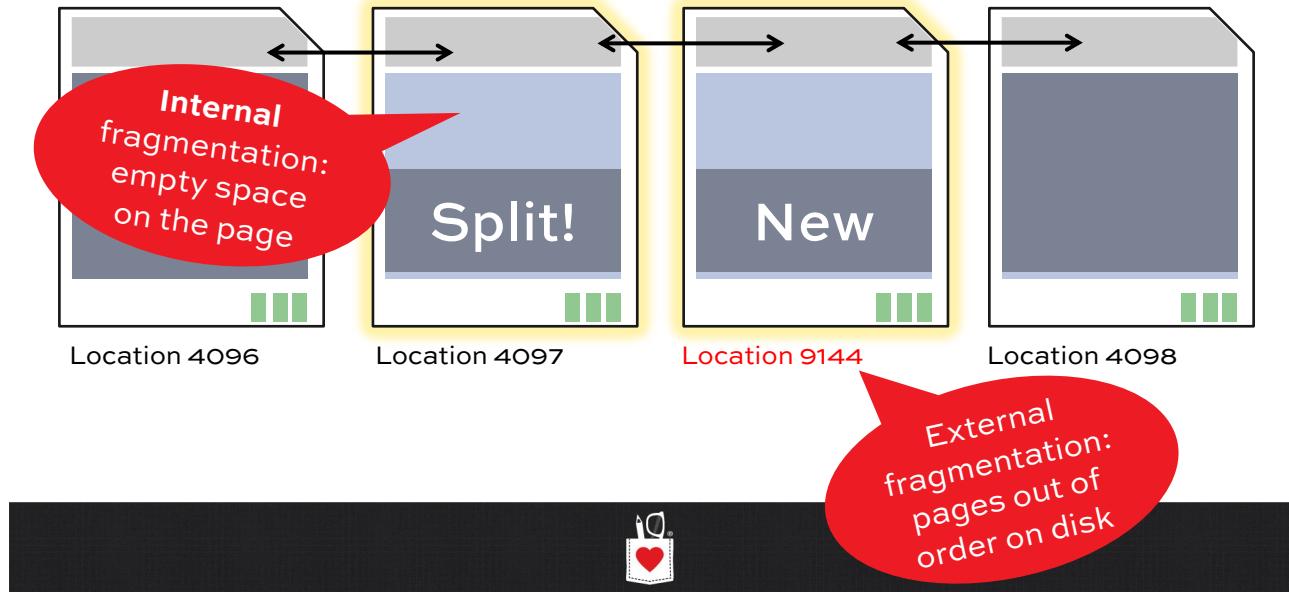
SQL allocates
another page.



But it won't be in
physical order.



There are two kinds of fragmentation here.



Monitoring this stuff

To track page splits:

Perfmon counter SQL Server:Access Methods – Page Splits/sec

To track external fragmentation:

sys.dm_db_index_physical_stats – avg_fragmentation_in_percent

To track internal fragmentation:

sys.dm_db_index_physical_stats – avg_page_space_used_in_percent





Great! I'll monitor
for page splits.



What does that actually mean though?

To find out, let's:

1. Create a table
2. Check page splits
3. Insert one row
4. Check page splits again

```
1  DROP TABLE IF EXISTS dbo.Customers;
2  GO
3  CREATE TABLE dbo.Customers
4    (Id INT IDENTITY(1,1) PRIMARY KEY CLUSTERED,
5     CustomerName VARCHAR(100));
6  GO
7  SELECT * FROM sys.dm_os_performance_counters
8  WHERE counter_name LIKE '%Splits%';
9  GO
10 INSERT INTO dbo.Customers(CustomerName)
11   VALUES ('Brent Ozar')
12  GO
13 SELECT * FROM sys.dm_os_performance_counters
14 WHERE counter_name LIKE '%Splits%';
```



```

7  SELECT * FROM sys.dm_os_performance_counters
8  WHERE counter_name LIKE '%Splits%';
9  GO
10 INSERT INTO dbo.Customers(CustomerName)
11      VALUES ('Brent Ozar')
12 GO
13 SELECT * FROM sys.dm_os_performance_counters
14 WHERE counter_name LIKE '%Splits%';

```

Results Messages

object_name	counter_name	instance_name	cntr_value	cntr_type
SQLServer:Access Methods	Page Splits/sec		30	272696576

object_name	counter_name	instance_name	cntr_val	cntr_type
SQLServer:Access Methods	Page Splits/sec		31	272696576

Went up by 1.

Adding a new page is also called a **page split**.



Monitoring this stuff

To track page splits:

Perfmon counter SQL Server:Access Methods – Page Splits/sec

There are other more accurate ways to monitor page splits, but if this lying thing came as a surprise to you, buckle up. It's about to get worse.

Lies

To track external fragmentation:

`sys.dm_db_index_physical_stats – avg_fragmentation_in_percent`

To track internal fragmentation:

`sys.dm_db_index_physical_stats – avg_page_space_used_in_percent`



Is **external**
fragmentation bad?



Pop quiz: if your pages are out of order

Will your maintenance tasks take longer?
(Backups, index & stats maintenance, CHECKDB)

Will queries reading from RAM be slower?

Will queries reading from disk be slower?

- What if you're seeking?
- What if you're scanning?

How do you fix it?



Pop quiz: if your pages are out of order

Will your maintenance tasks take longer? No
(Backups, index & stats maintenance, CHECKDB)

Will queries reading from RAM be slower? No

Will queries reading from disk be slower?

- What if you're seeking? No
- What if you're scanning? Maybe,
but...

How do you fix it?



Systems Performance by Brendan Gregg

1 CPU cycle	0.3 ns
Level 1 cache access	0.9 ns
Level 2 cache access	2.8 ns
Level 3 cache access	12.9 ns
Main memory access	120 ns
Solid-state disk I/O	50-150 µs
Rotational disk I/O	1-10 ms
Internet: SF to NYC	40 ms
Internet: SF to UK	81 ms
Internet: SF to Australia	183 ms



Systems Performance by Brendan Gregg

1 CPU cycle	0.3 ns	1 s
Level 1 cache access	0.9 ns	3 s
Level 2 cache access	2.8 ns	9 s
Level 3 cache access	12.9 ns	43 s
Main memory access	120 ns	6 min
Solid-state disk I/O	50-150 µs	2-6 days
Rotational disk I/O	1-10 ms	1-12 months
Internet: SF to NYC	40 ms	4 years
Internet: SF to UK	81 ms	8 years
Internet: SF to Australia	183 ms	19 years



Systems Performance by Brendan Gregg

1 CPU cycle	0.3 ns	1 s
Level 1 cache access	0.9 ns	3 s
Level 2 cache access	2.8 ns	9 s
Level 3 cache access	12.9 ns	43 s
Main memory access	120 ns	6 min
Solid-state disk I/O	20 µs	2-6 days
Rotational disk I/O	1 ms	1-12 months
Internet: SF to NYC	1 ms	4 years
Internet: SF to UK	81 ms	8 years
Internet: SF to Australia	183 ms	19 years



They don't notice the difference down here.

Users want this.

Pop quiz: if your pages are out of order

Will your maintenance tasks take longer? **No**
(Backups, index & stats maintenance, CHECKDB)

Will queries reading from RAM be slower? **No**

Will queries reading from disk be slower?

- What if you're seeking? **No**
- What if you're scanning? **Maybe, but...**

How do you fix it?

This sucks.



Fixing out-of-order pages

	Reorganize	Rebuild
What does it do?	Shuffles individual pages around, one at a time	Builds a whole new copy of the index
Are all pages in physical order when we're done?	No	Maybe, but only if you use MAXDOP 1
Can it be done online?	Yes	Enterprise Edition only
Can it be interrupted?	Yes, and your progress is kept	No
Is it a logged operation?	Yes	Yes
Does it slow down mirrors, AGs, log shipping?	Yes	Yes



Users hate
your “fix.”



Pop quiz: if your pages are out of order

Will your maintenance tasks take longer?

(Backups, index & stats maintenance, CHECKDB)

Will queries reading from RAM be slower?

Defragmenting makes this WORSE

Doesn't help

Will queries reading from disk be slower?

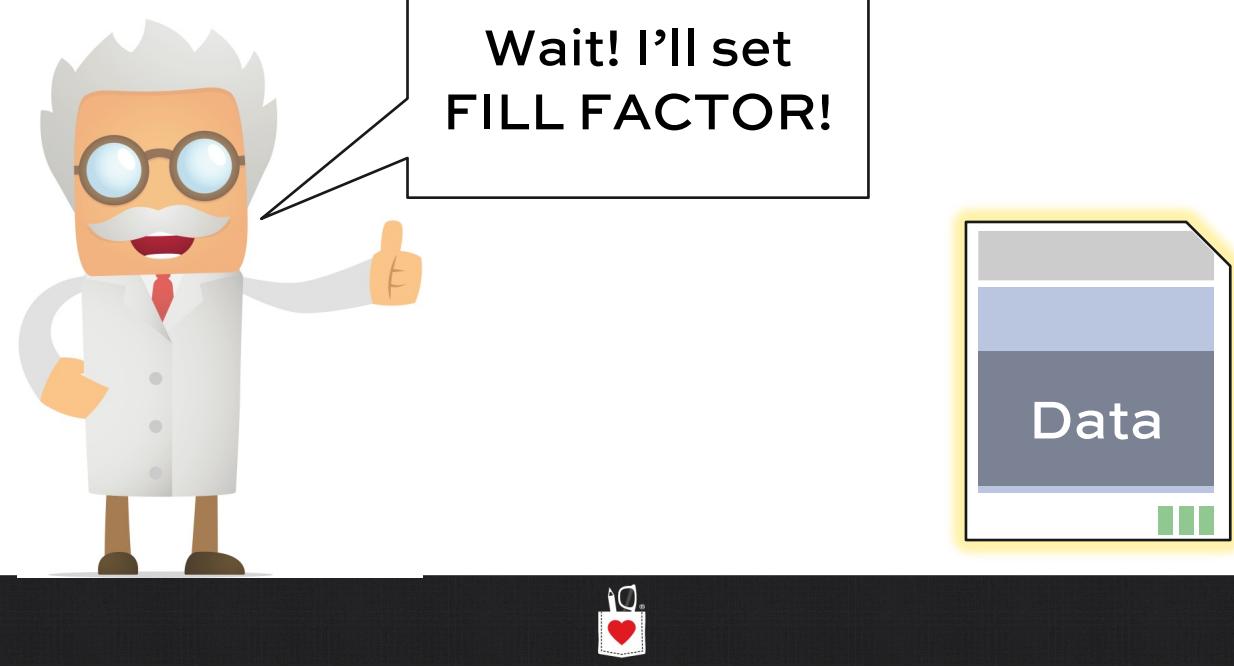
Doesn't help ?

- What if you're seeking?
- What if you're scanning?

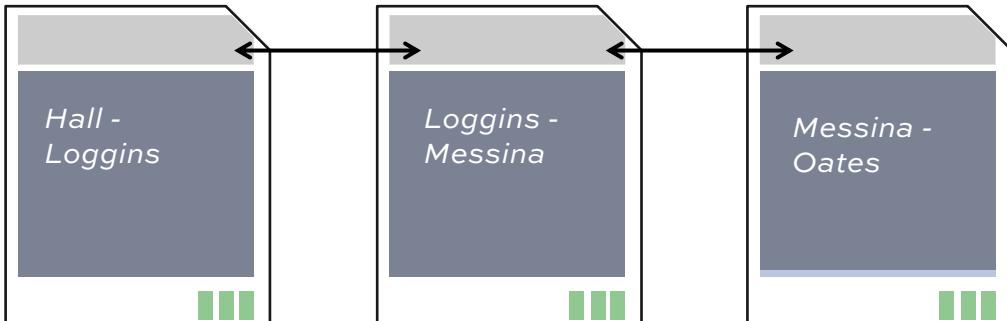
Helps – but not much

How do you fix it?

I'm not so sure we should, but if we do, it shouldn't be with constant index shuffling.



McDonald, Michael moves to our town.



He needs to go here, but there's no empty space.

But what if there was?



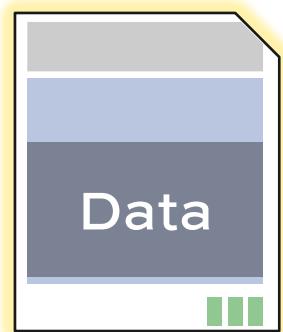
Fill factor leaves empty space on the page.

Default set at the server level: 100% (and 0 means the same thing)

So the idea is:

- Set fill factor to something lower (like 70-80)
- Leave empty space on every page
- When new records are added, we'll have space so we don't need to split pages

Let's see a couple examples.



dbo.Users - Clustered Index

dbo.Users Clustered index on ID

Users are ordered by ID, an identity field

Do you want empty space on this page?

- Inserts?
- Updates?
- Deletes?

What's the right fill factor?

ID	Rep	CreationDate	DisplayName	LastAccessDate	Location	Age	AboutMe	
1	2406	7/12/09 10:51 PM	Jeff Atwood	4/1/10 10:35 AM	El Cerrito, CA	39	Software Developer-<	29 <2>Software Developer-<
558	953	7/15/09 9:46 AM	marco.ragagna	2/18/10 7:14 AM	AU	39 Hey, at least I know 6502!	39 Hey, at least I know 6502!	
559	593	7/15/09 9:46 AM	moobaa	3/4/10 6:47 PM	United Kingdom	30 <3>Tech geek and hopeless	30 <3>Tech geek and hopeless	
560	83	7/15/09 9:46 AM	rjstelling	3/16/10 8:45 AM	London	NULL	NULL	
561	716	7/15/09 9:47 AM	balph4	10/29/09 2:11 PM	Bremen, Germany	24 NULL	24 NULL	
562	43	7/15/09 9:47 AM	Mike McClelland	4/1/10 6:59 AM	UK	30 <p>I'm a Software Enginee	30 <p>I'm a Software Enginee	
563	101	7/15/09 9:48 AM	arturh	3/5/10 12:13 AM	Kirkland, WA	37 <p>Father, SDE@Microsoft .	37 <p>Father, SDE@Microsoft .	

From <http://StackOverflow.com's Creative Commons Data Dump>

dbo.Users - IX_LastAccessDate

dbo.Users non-clustered on LastAccessDate

Users are ordered by the last date they visited Stack Overflow

Do you want empty space on this page?

- Inserts?
- Updates?
- Deletes?

What's the right fill factor?

LastAccessDate	ID	LastAccessDate	ID	LastAccessDate	ID	LastAccessDate	ID
7/31/08 12:00 AM	-1	7/15/09 8:53 AM	445	7/15/09 9:10 PM	200	8/11/09 7:17 PM	39
7/15/09 7:08 AM	22	7/15/09 8:58 AM	457	7/16/09 6:22 AM	678	8/12/09 2:54 PM	943
7/15/09 7:10 AM	33	7/15/09 9:17 AM	501	7/17/09 2:30 AM	131	8/13/09 4:26 PM	364
7/15/09 7:11 AM	40	7/15/09 9:28 AM	524	7/17/09 9:30 AM	297	8/15/09 5:03 PM	910
7/15/09 7:11 AM	41	7/15/09 9:30 AM	527	7/17/09 8:43 PM	998	8/17/09 8:42 AM	202
7/15/09 7:11 AM	44	7/15/09 9:58 AM	587	7/18/09 12:38 PM	394	8/17/09 10:11 AM	628
7/15/09 7:12 AM	52	7/15/09 10:00 AM	594	7/18/09 2:15 PM	924	8/17/09 10:33 AM	157
7/15/09 7:13 AM	64	7/15/09 10:02 AM	597	7/19/09 10:26 PM	336	8/17/09 4:24 PM	1006
7/15/09 7:13 AM	65	7/15/09 10:21 AM	618	7/20/09 1:06 PM	849	8/18/09 8:06 AM	511
7/15/09 7:14 AM	68	7/15/09 10:25 AM	347	7/21/09 7:22 AM	881	8/18/09 9:00 AM	262
7/15/09 7:15 AM	73	7/15/09 10:26 AM	623	7/23/09 11:53 AM	503	8/18/09 9:43 AM	210
7/15/09 7:17 AM	87	7/15/09 10:28 AM	629	7/23/09 12:56 PM	446	8/18/09 10:22 AM	673
7/15/09 7:18 AM	92	7/15/09 10:32 AM	638	7/24/09 12:15 AM	407	8/18/09 1:05 PM	959
7/15/09 7:20 AM	110	7/15/09 10:36 AM	642	7/24/09 12:08 PM	488	8/18/09 1:14 PM	643
7/15/09 7:26 AM	139	7/15/09 10:46 AM	661	7/24/09 3:17 PM	144	8/18/09 1:32 PM	525
7/15/09 7:36 AM	173	7/15/09 10:59 AM	686	7/26/09 8:48 AM	615	8/18/09 1:38 PM	159
7/15/09 7:38 AM	179	7/15/09 11:14 AM	722	7/27/09 2:27 PM	337	8/18/09 2:28 PM	690
7/15/09 7:46 AM	223	7/15/09 11:26 AM	756	7/28/09 3:17 AM	174	8/19/09 9:22 AM	750
7/15/09 7:57 AM	258	7/15/09 11:48 AM	803	7/28/09 12:34 PM	118	8/19/09 2:46 PM	887
7/15/09 7:58 AM	259	7/15/09 11:50 AM	681	7/28/09 9:28 PM	201	8/20/09 3:45 AM	749
7/15/09 8:03 AM	279	7/15/09 11:51 AM	808	7/28/09 10:53 PM	656	8/20/09 8:29 AM	268
7/15/09 8:06 AM	286	7/15/09 12:00 PM	822	7/29/09 12:06 PM	751	8/20/09 10:32 AM	901
7/15/09 8:12 AM	305	7/15/09 12:07 PM	841	7/30/09 12:26 PM	754	8/20/09 12:05 PM	752
7/15/09 8:13 AM	313	7/15/09 12:07 PM	842	7/31/09 6:36 PM	11	8/20/09 3:03 PM	272
7/15/09 8:14 AM	320	7/15/09 12:18 PM	878	8/1/09 12:09 PM	775	8/21/09 7:29 PM	487
7/15/09 8:16 AM	323	7/15/09 12:18 PM	779	8/1/09 7:29 PM	670	8/22/09 11:58 AM	580
7/15/09 8:25 AM	349	7/15/09 12:47 PM	973	8/1/09 11:33 PM	85	8/24/09 11:22 AM	56
7/15/09 8:26 AM	354	7/15/09 12:48 PM	819	8/3/09 10:13 PM	300	8/24/09 11:33 PM	889
7/15/09 8:27 AM	97	7/15/09 12:52 PM	987	8/5/09 12:06 PM	918	8/25/09 8:57 PM	150
7/15/09 8:27 AM	363	7/15/09 12:54 PM	992	8/6/09 5:46 AM	726	8/26/09 2:43 AM	851
7/15/09 8:33 AM	384	7/15/09 1:28 PM	931	8/6/09 10:42 AM	370	8/26/09 2:53 AM	117
7/15/09 8:37 AM	402	7/15/09 2:19 PM	728	8/6/09 12:37 PM	398	8/26/09 9:59 AM	546
7/15/09 8:47 AM	431	7/15/09 2:20 PM	217	8/9/09 5:28 AM	45	8/27/09 3:54 AM	239
7/15/09 8:52 AM	440	7/15/09 4:23 PM	975	8/10/09 12:25 AM	389	8/27/09 2:10 PM	271

From <http://StackOverflow.com's Creative Commons Data Dump>



The default fill factor
of 100% keeps your
database small.



When you set fill factor
 $<> 100\%$, guess what
you're really doing.





When you set fill factor
 $<> 100\%$, guess what
you're really doing.

That's right: setting fill factor $<> 100\%$ actually causes...

Internal fragmentation

Pop quiz: if your pages have empty space

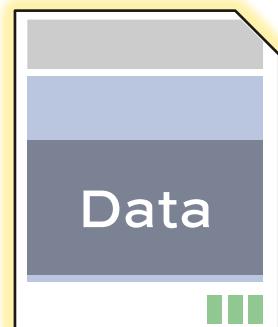
Will your maintenance tasks take longer?
(Backups, index & stats maintenance, CHECKDB)

Will queries reading from RAM be slower?

Will queries reading from disk be slower?

- What if you're seeking?
- What if you're scanning?

How do you fix it?



Pop quiz: if your pages have empty space

Will your maintenance tasks take longer? **YES**
(Backups, index & stats maintenance, CHECKDB)

Will queries reading from RAM be slower? **YES**

Will queries reading from disk be slower?

- What if you're seeking? **No**
- What if you're scanning? **YES**



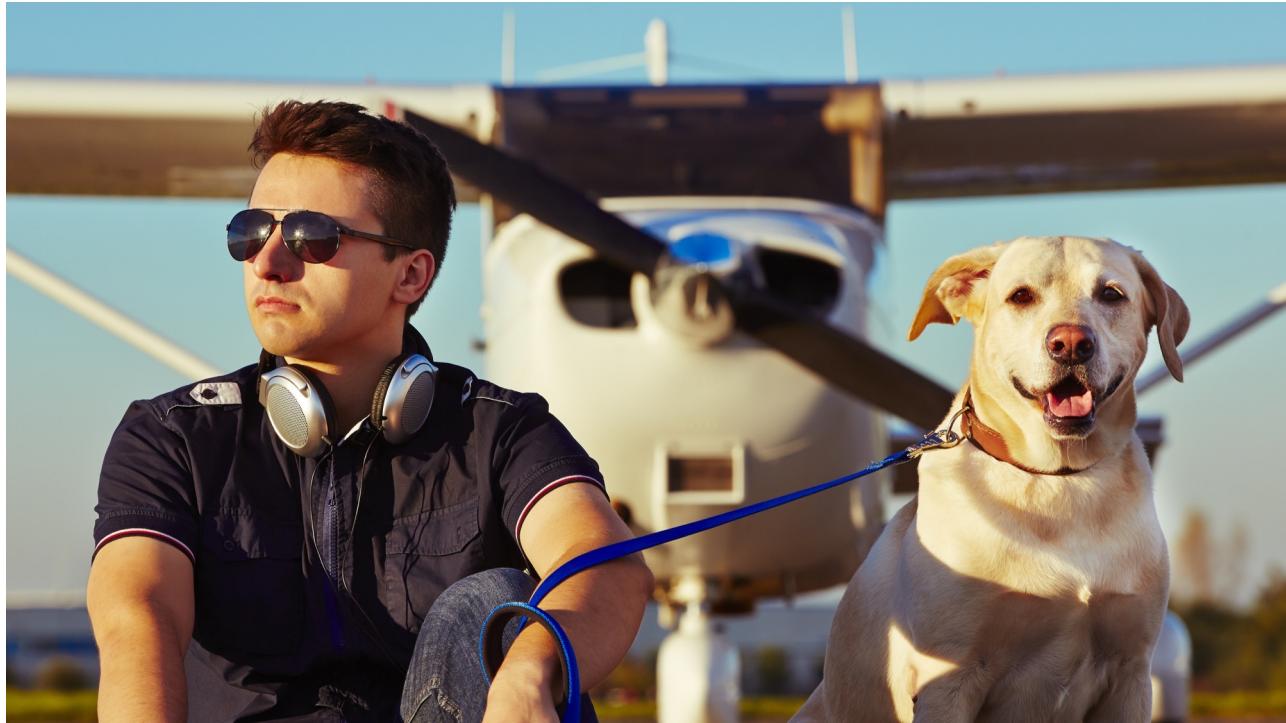
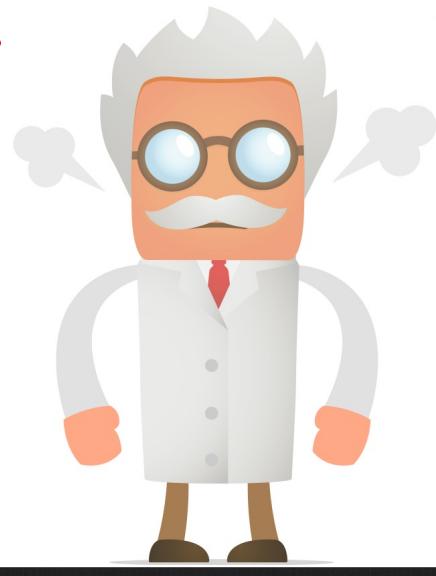
THE PROBLEM IS US.

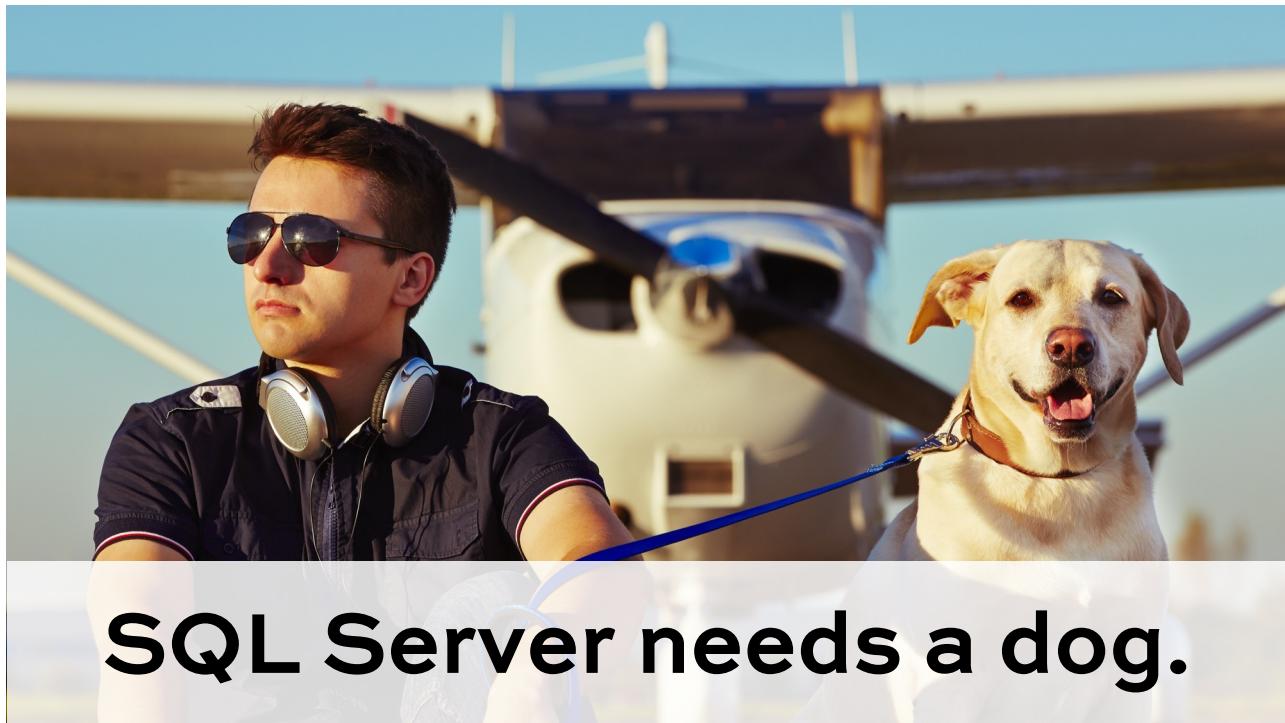
External fragmentation doesn't really matter

But to “fix” it, we’ve been:

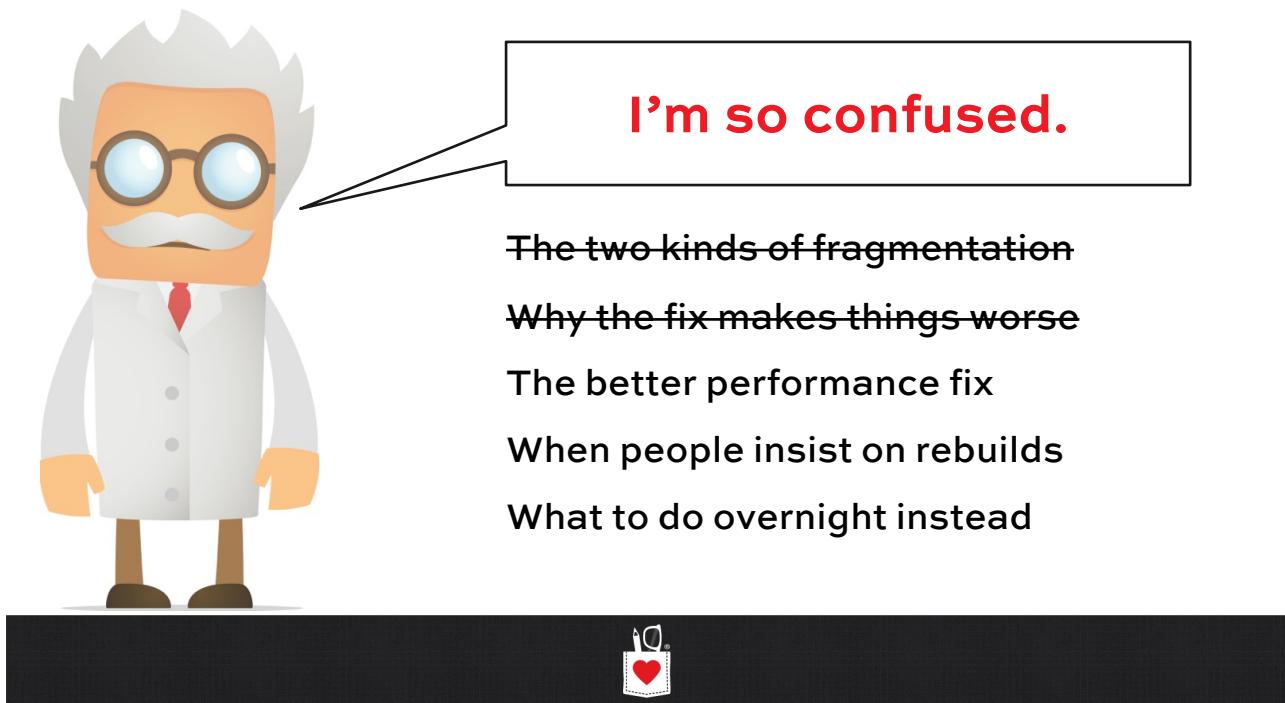
- Rebuilding/reorging our indexes daily, which actually **makes some things worse**
- Setting fill factor, which **causes internal fragmentation**

And internal fragmentation **DOES** matter
(and we’ve been the ones causing it all along!)





SQL Server needs a dog.



The short answer

1. Set fill factor at the server level back to 100%.
2. Use sp_Blitz and sp_BlitzIndex to check for index fill factors <80%, and set those back up to 80% or higher (or 100%).
3. Rebuild your indexes once to atone for those past sins, getting them back up to 100% fill factor.
4. Monitor the right metric, and use the scientific method to improve it.



How SQL Server Schedules CPU

What's Running Now

What's Waiting (Queue)



How SQL Server Schedules CPU

What's Running Now

`SELECT * FROM
dbo.Restaurants
(By Brent)`

What's Waiting (Queue)



How SQL Server Schedules CPU

What's Running Now

`SELECT * FROM
dbo.Restaurants
(By Brent)`

What's Waiting (Queue)

`SELECT * FROM
dbo.Tattoos
(By Jeremiah)`

`SELECT * FROM
dbo.Rabbits
(By Kendra)`

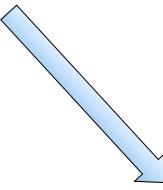


How SQL Server Schedules CPU

What's Running Now
`SELECT * FROM
dbo.Restaurants
(By Brent)`

What's Waiting (Queue)
`SELECT * FROM
dbo.Tattoos
(By Jeremiah)`

`SELECT * FROM
dbo.Rabbits
(By Kendra)`



How SQL Server Schedules CPU

What's Running Now

What's Waiting (Queue)



`SELECT * FROM
dbo.Tattoos
(By Jeremiah)`

`SELECT * FROM
dbo.Rabbits
(By Kendra)`

`SELECT * FROM
dbo.Restaurants
(By Brent)`



SQL Server Waits For:

Resources:

CPU, memory, storage, network, latches, locks

Stuff outside of SQL Server (Preemptive):

COM, OLEDB, CLR

System Tasks:

Lazywriter, trace, full text search



Is this SQL Server working hard?

What's Running Now

SELECT * FROM
dbo.Restaurants
(By Brent)

What's Waiting (Queue)



What about now?

What's Running Now
SELECT * FROM
dbo.Restaurants
(By Brent)

What's Waiting (Queue)
SELECT * FROM dbo.Tattoos
(By Jeremiah)
SELECT * FROM dbo.Rabbits
(By Kendra)



Or now?

What's Running Now
SELECT * FROM
dbo.Restaurants
(By Brent)

What's Waiting (Queue)
SELECT * FROM dbo.Tattoos
(By Jeremiah)
SELECT * FROM dbo.Rabbits
(By Kendra)

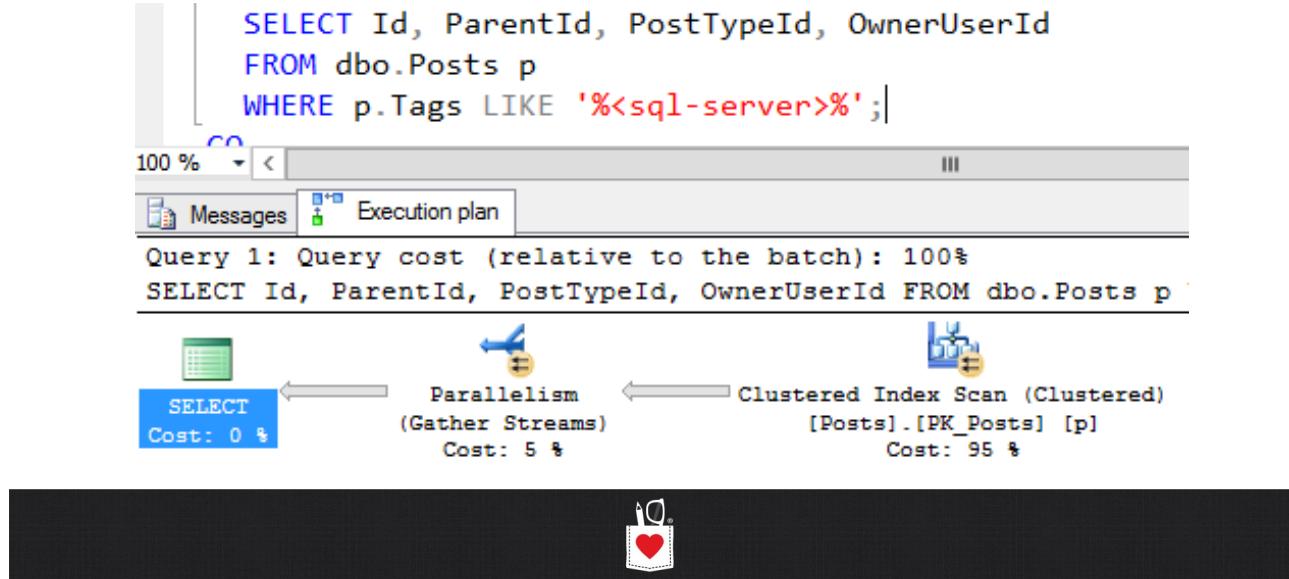
EXEC sp_WhoIsActive
(By Adam Machanic)

BACKUP DATABASE SQLbits
(By Ola Hallengren)

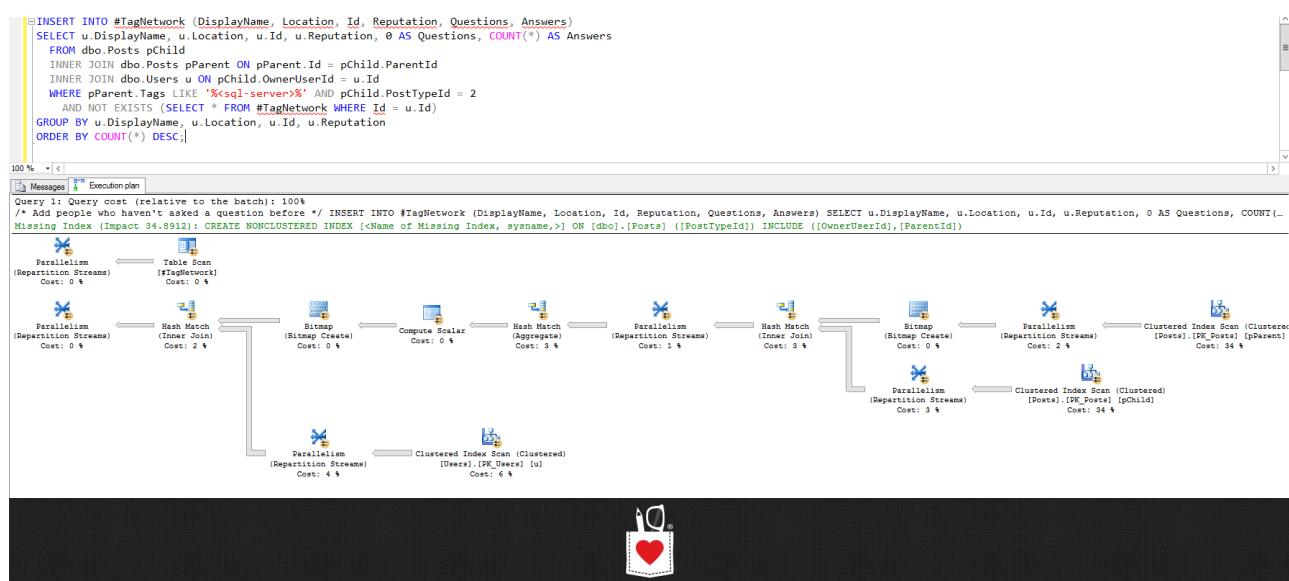
DBCC CHECKDB
(By Paul Randal)



Some queries are simple



Some queries have a lot going on



These were simple little queries.

What's Running Now
SELECT * FROM
dbo.Restaurants
(By Brent)

What's Waiting (Queue)
SELECT * FROM
dbo.Tattoos
(By Jeremiah)

SELECT * FROM
dbo.Rabbits
(By Kendra)



But one query looks more like this:

What's Running Now
Index scan of Table1

What's Waiting (Queue)
Index scan of Table2

Index scan of Table3

Index scan of Table4



So for each core, you'll likely see:

What's Running Now

**One task running,
using CPU**

What's Waiting (Queue)

**Several
(or several DOZEN)
tasks stacked up
waiting on stuff**



**If we do this for one second,
how many seconds of waits?**

What's Running Now

**One task running,
using CPU**

What's Waiting (Queue)

**15 tasks waiting on
storage**

**3 tasks waiting on
locks**

1 task waiting on CPU



Hours of wait time per hour



How hard is your server working?

Dynamic Management View (DMV)
sys.dm_os_wait_stats

Tracked cumulatively over time

Trend on an hourly basis and break it out by:

- Weekday vs weekend
- Business hours vs after hours
- Maintenance windows
(backups, DB maintenance)



Wait Time per core per hour

Near zero – Your server isn't doing anything.
(Although that doesn't mean every query is fast.)

1 hour of waits – You're still not doing much.

Multiple hours per core – now we're working!
And we should probably be tuning.



Free script: sp_BlitzFirst

Totally free open source diagnostic tool.

Installs in the master by default, but can go anywhere.

Runs in 5 seconds ideally, but can be more under load.

By default, shows waits for a 5-second sample now.

@SinceStartup = 1 shows waits since, uh, startup.



Solve the emergencies, then focus on waits, in order.

Some waits are like poison:
any occurrences of them have to be fixed first.

After that, work through the list.

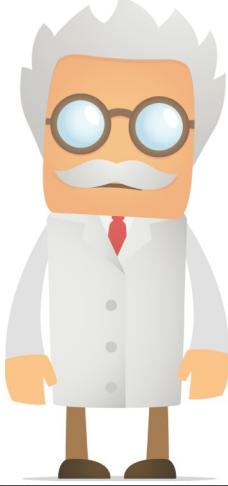
You'll never make them go away,
but you have to focus on them in order.

	Priority	FindingsGroup	Finding	URL	Details	How
1	0	sp_BlitzFirst 2017-03-0...	From Your Community Volunteers	http://FirstResponderKit.org/	Click To See Details -- We hope you found this t...	NUL
2	10	Server Performance	Poison Wait Detected: THREADPOOL	http://www.brentozar.com/go/poi...	Click To See Details -- For 0 seconds over the I...	?CI
3	50	Server Performance	High CPU Utilization	http://www.BrentOzar.com/go/cpu	Click To See Details -- 98%. Ring buffer details: ...	NUL
4	200	Wait Stats	CXPACKET	http://www.brentozar.com/sql/wai...	Click To See Details -- For 25750 seconds over ...	?CI
5	200	Wait Stats	SOS_SCHEDULER_YIELD	http://www.brentozar.com/sql/wai...	Click To See Details -- For 237 seconds over th...	?CI
6	200	Wait Stats	MEMORY_ALLOCATION_EXT	http://www.brentozar.com/sql/wai...	Click To See Details -- For 55 seconds over the ...	?CI

Wait stats decoder ring

Wait type	Means	sp_BlitzCache @SortOrder =
CXPACKET, CXCONSUMER	Parallelism	Reads, CPU
LCK*	Blocking locks	Duration, reads
PAGEIOLATCH	Reading data from data files	Reads
SOS_SCHEDULER_YIELD	Waiting for CPU time	CPU
RESOURCE_SEMAPHORE	Query needs memory to start	Unused grant
WRITELOG	Writing to the transaction log	Writes
HADR_SYNC_COMMIT	Writing to AG sync secondaries	Writes





What if someone wants to play doctor?

~~The two kinds of fragmentation~~

~~Why the fix makes things worse~~

~~The better performance fix~~

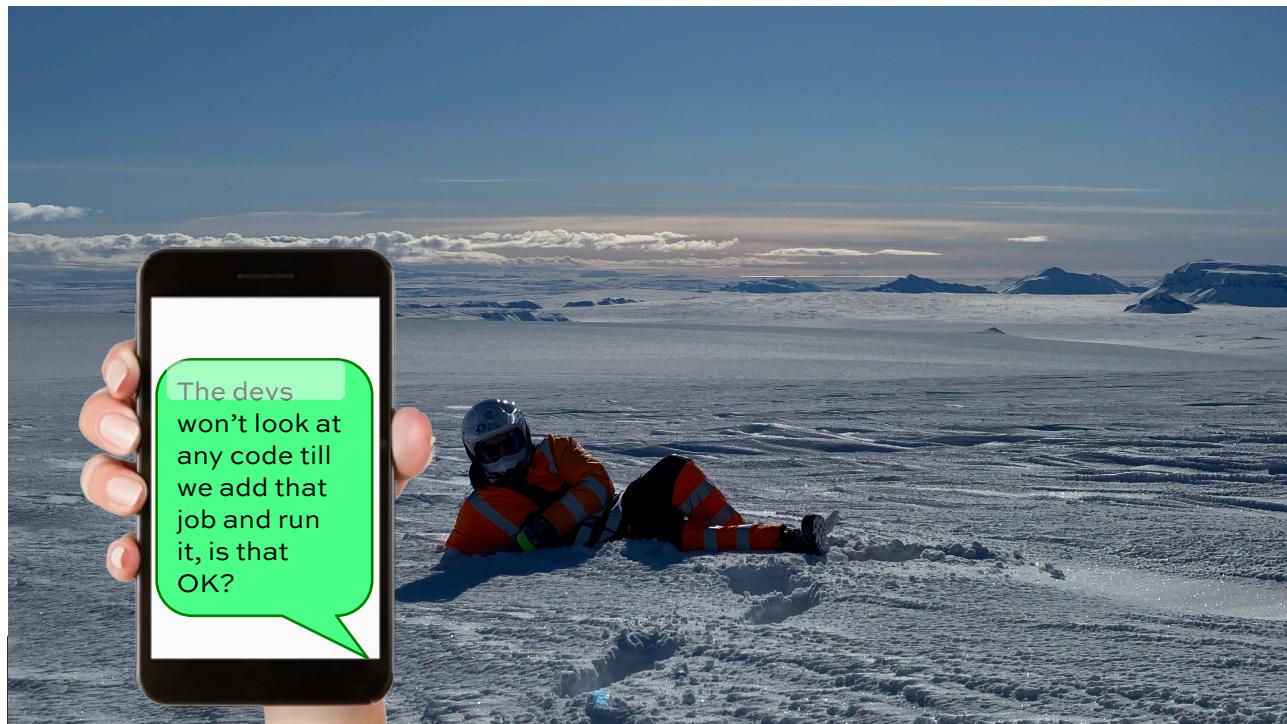
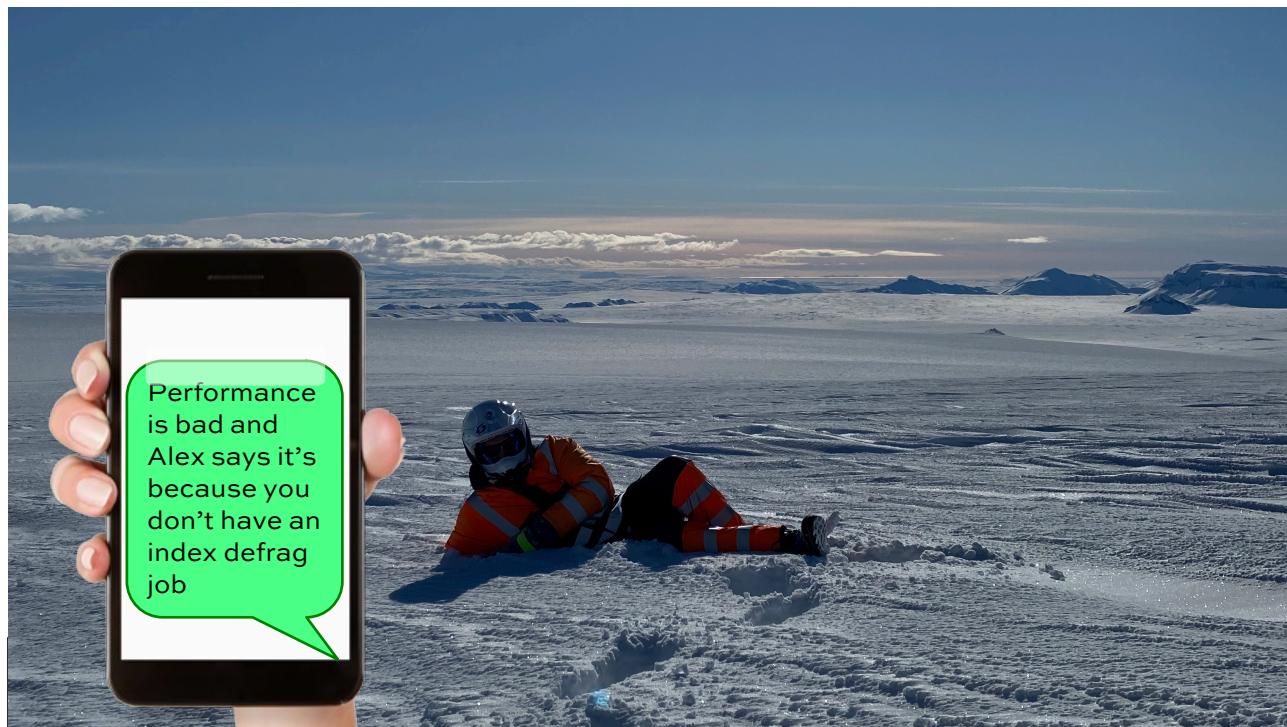
When people insist on rebuilds

What to do overnight instead



You're on vacation when suddenly....





How to react

“The problem is high fragmentation.”

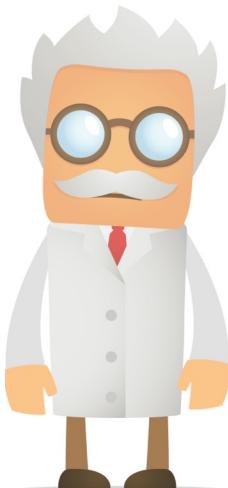
“So you guarantee that rebuilding the indexes will fix it?”

“Uhhhh...”

“What metric shall we monitor to know that the rebuild fixed it?”

“Ummmm...query runtime?”

“Great. Give me a query, and the current average runtime, and what you expect it’ll be afterwards.”



So what do I do daily?

~~The two kinds of fragmentation~~

~~Why the fix makes things worse~~

~~The better performance fix~~

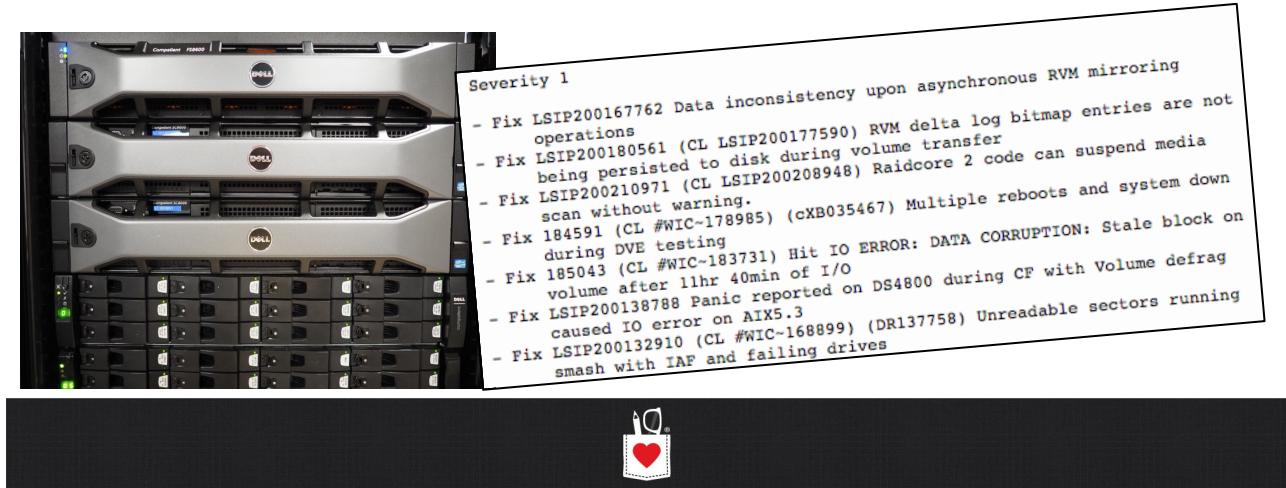
~~When people insist on rebuilds~~

~~What to do overnight instead~~

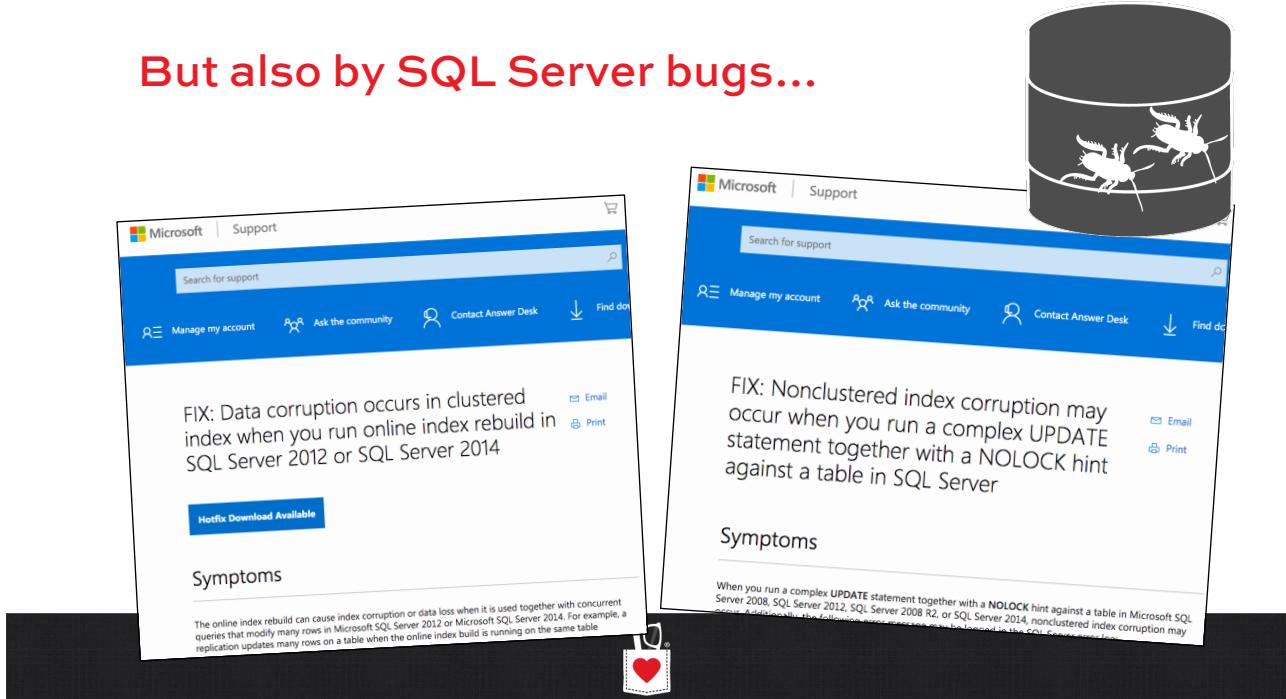


People get fired for losing data.

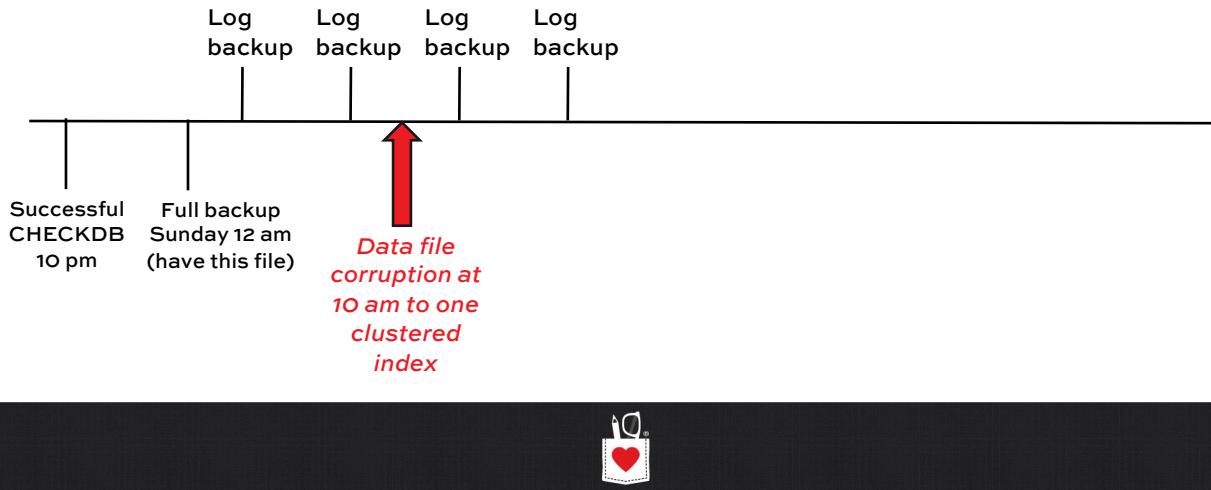
Data corruption is sometimes caused by storage failures and bugs...



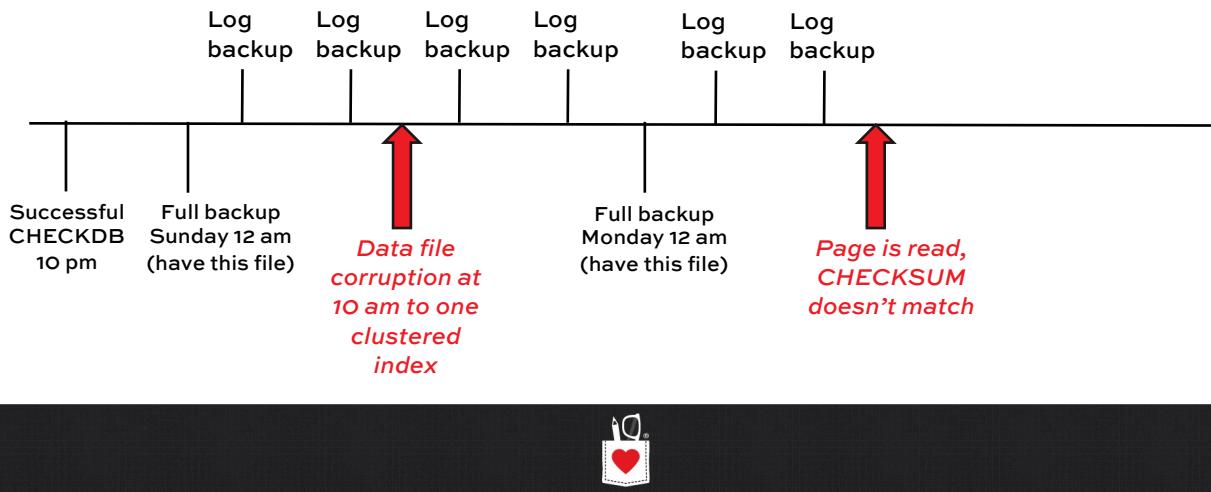
But also by SQL Server bugs...



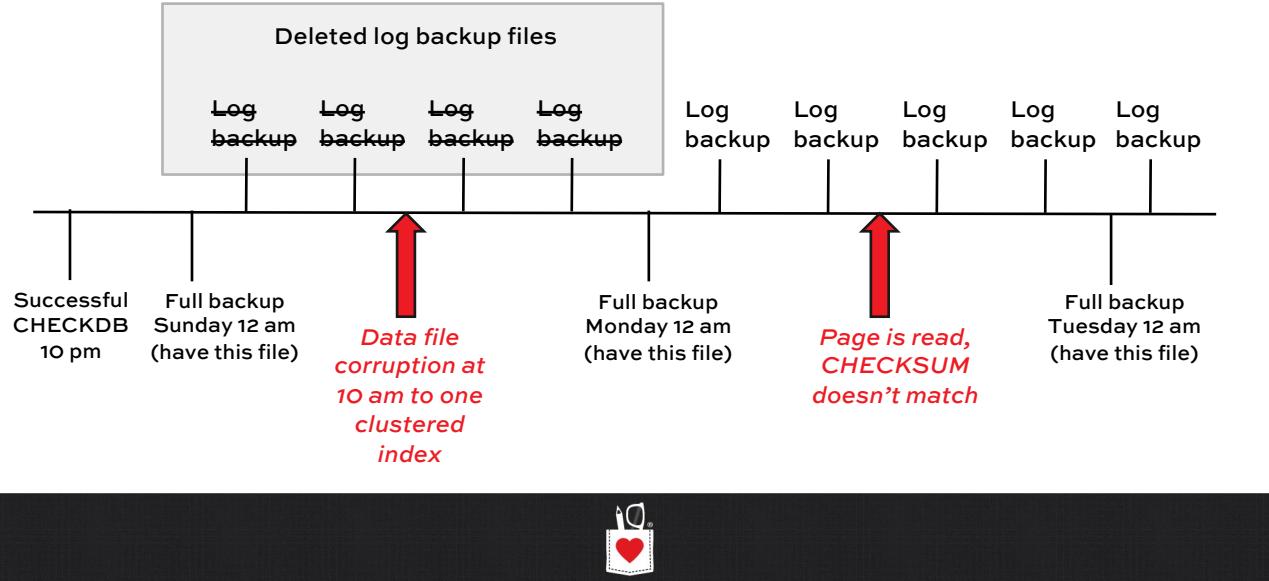
We get corruption sometime on Sunday



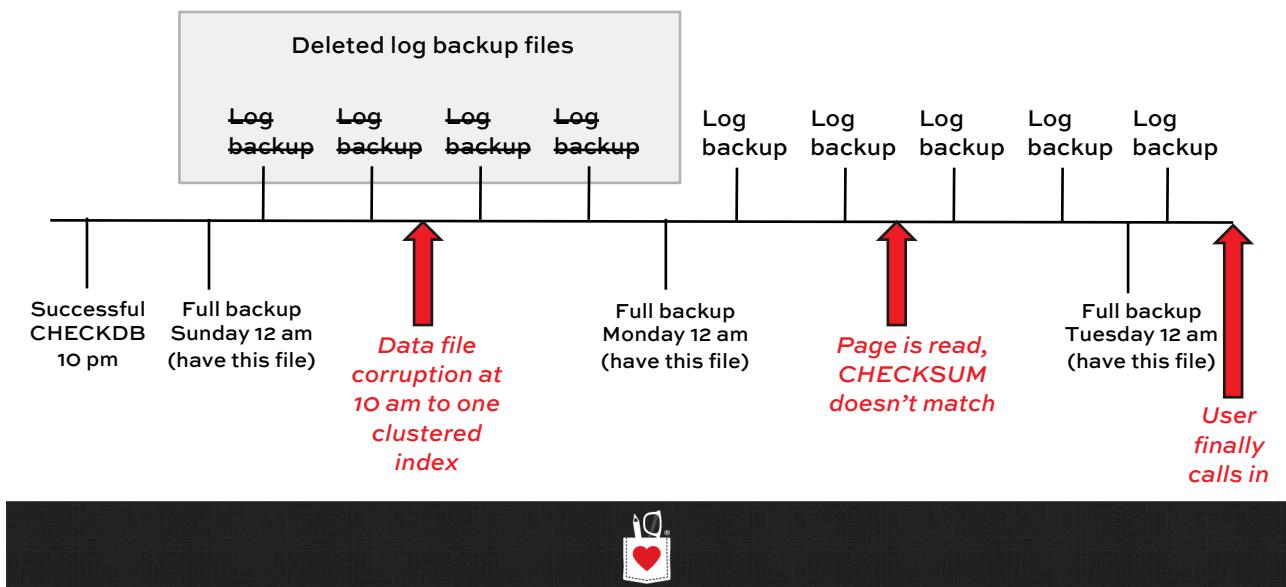
Corruption is detected – but not acted on



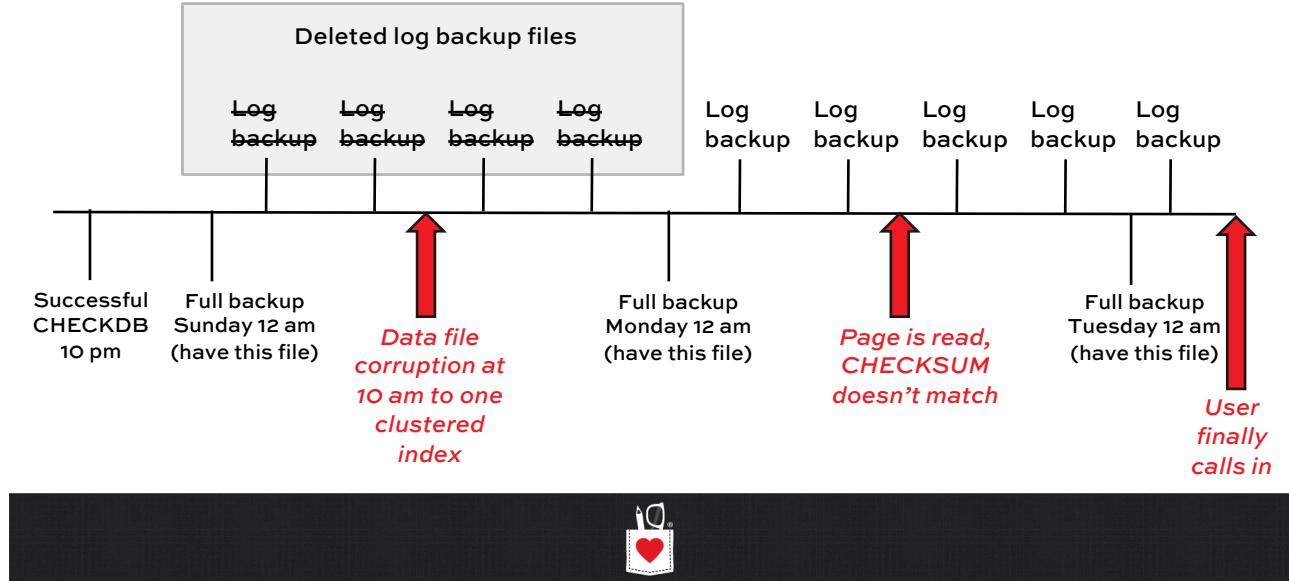
We only keep 24 hours of log backup files



Will you lose data? How much?



What could you do to lose less data?



Ways to lose less data

Set up alerts for corruption errors: <https://BrentOzar.com/go/alerts>

Keep transaction log backup files longer
(dictated by your RPO, RTO, and CHECKDB frequency)

Run CHECKDB more often (and act fast on the outputs)

The more of these you do, the better your chances are to keep your job.



**DBAs don't get fired
for slow performance.**



**DBAs get fired
for losing data.**



Run CHECKDB as frequently as practical.

Forget shuffling indexes around every night.

It's time to run CHECKDB in that time window instead.



Yay, it's over!

Recap



What you learned

External fragmentation = pages out of order on disk
(Only matters when you're reading from disk)

Internal fragmentation = empty space on pages
(Matters a lot, and fix this with judicious rebuilds)

Fill factor = internal fragmentation
(Set this, and you're making internal fragmentation worse)

**Index maintenance is fine,
but it doesn't solve most wait types.**
Find your top wait type with sp_BlitzFirst and focus on it.



What you'll do tomorrow

1. Set fill factor at the server level back to 100%.
2. Use sp_Blitz and sp_BlitzIndex to check for index fill factors <80%, and set those back up to 80% or higher (or 100%).
3. Rebuild your indexes once to atone for those past sins, getting them back up to 100% fill factor.
4. Monitor wait time per hour, and use the scientific method to improve it.

